



# OPENSIFT

OpenShift Installation using Ansible  
& Terraform

**DOCUMENT VERSION V 1.0**

## Authors

**Adam Muganwa**  
Linux Systems Engineer  
[adam.muganwa@ibm.com](mailto:adam.muganwa@ibm.com)  
Slack: Adam Muganwa

**Ben Rossouw**  
IT Specialist & Infrastructure Manager  
[brossouw@za.ibm.com](mailto:brossouw@za.ibm.com)  
Slack: Ben Rossouw

# Openshift 4x Installation vSphere

## Prerequisites

Define OCP cluster id (cluster name) I will be using **envisage**

### User-provisioned DNS requirements

[https://docs.openshift.com/container-platform/4.2/installing/installing\\_vsphere/installing-vsphere.html#installation-dns-user-infra\\_installing-vsphere](https://docs.openshift.com/container-platform/4.2/installing/installing_vsphere/installing-vsphere.html#installation-dns-user-infra_installing-vsphere)

## Setup Bastion node on the network.

Bastion Node used in this documentation is a Centos VM

### Cleaning up previous installation

#### In installation folder

```
rm -rf*: rm -rf auth -R *.ign .openshift_install* metadata.json
```

#### In vsphere folder

```
rm -rf terraform.tfstate .terraform/ -R
```

### Installing required packages

```
yum install ansible -y  
yum install git -y  
yum install wget unzip -y  
yum install bind-utils -y  
yum install vim -y
```

### Check DNS entries

```
nslookup api.envisage.clientcenter.local
```

Should return the name and ip address

```
[root@Bastionhost ~]# nslookup api.envisage.clientcenter.local  
Server:      172.30.12.50  
Address:     172.30.12.50#53  
  
Name:   api.envisage.clientcenter.local  
Address: 172.30.50.110
```

### Check the rest of the records.

```
nslookup api-int.envisage.clientcenter.local  
nslookup *.apps.envisage.clientcenter.local  
nslookup bootstrap-0.envisage.clientcenter.local  
nslookup control-plane-0.envisage.clientcenter.local  
nslookup control-plane-1.envisage.clientcenter.local  
nslookup control-plane-2.envisage.clientcenter.local  
nslookup compute-0.envisage.clientcenter.local  
nslookup compute-1.envisage.clientcenter.local  
nslookup compute-2.envisage.clientcenter.local
```

```
nslookup compute-3.envisage.clientcenter.local  
nslookup compute-4.envisage.clientcenter.local
```

```
nslookup -type=SRV _etcd-server-ssl._tcp.envisage.clientcenter.local
```

```
[root@Bastionhost ~]# nslookup -type=SRV _etcd-server-ssl._tcp.envisage.clientcenter.local  
Server: 172.30.12.50  
Address: 172.30.12.50#53  
  
_etcd-server-ssl._tcp.envisage.clientcenter.local service = 0 10 2380 etcd-2.envisage.clientcenter.local.  
_etcd-server-ssl._tcp.envisage.clientcenter.local service = 0 10 2380 etcd-1.envisage.clientcenter.local.  
_etcd-server-ssl._tcp.envisage.clientcenter.local service = 0 10 2380 etcd-0.envisage.clientcenter.local.
```

## Setup Load Balancer

### Download installer files

```
git clone https://github.com/adamsocool/Openshift4x-Installation
```

### copy the downloaded [loadbalancer.yml](#)

```
sudo cp loadbalancer.yml /etc/ansible
```

### Set up Load balancer/http server using playbook

```
cd /etc/ansible
```

### edit hosts file and add your load balancer to the file

```
vim hosts
```

```
[loadbalancer]  
172.30.50.110  
~
```

### generate ssh-key and copy to loadbalancer

```
ssh-keygen
```

```
ssh-copy-id root@172.50.30.110
```

### Check that you can connect to the server using ansible

```
ansible -m ping loadbalancer -o
```

```
[root@DevOps ansible]# ansible -m ping loadbalancer -o  
172.30.50.110 | SUCCESS => {"ansible_facts": {"discovered_interpreter_python": "/usr/libexec/platform-python"}, "changed": false, "ping": "pong"}  
[root@DevOps ansible]#
```

Run the playbook

ansible-playbook loadbalancer.yml

```
[root@DevOps ansible]# ansible-playbook loadbalancer.yml

PLAY [setup loadbalancer] ****
TASK [Gathering Facts] ****
ok: [172.30.50.110]

TASK [disable firewalld] ****
ok: [172.30.50.110]

TASK [Ensure SELinux is set to enforcing mode] ****
ok: [172.30.50.110]

TASK [Install a list of packages] ****
ok: [172.30.50.110]

TASK [turn on haproxy_connect_any] ****
changed: [172.30.50.110]

TASK [start and enable haproxy] ****
ok: [172.30.50.110]

TASK [start and enable httpd] ****
ok: [172.30.50.110]

TASK [bind httpd to port 8080] ****
ok: [172.30.50.110]

TASK [restart haproxy] ****
changed: [172.30.50.110]

PLAY RECAP ****
172.30.50.110 : ok=9    changed=2    unreachable=0    failed=0    skipped=0
                  rescued=0   ignored=0
```

## Login to Load Balancer

This configuration will add the following load balancer entries to HA-Proxy:

openshift-api-server (*port 6443*)

machine-config-server (*port 22623*)

ingress-http (*port 80*)

ingress-https (*port 443*)

Add the following configuration lines to /etc/haproxy/haproxy.cfg:

(*replace the IP addresses!*)

```
cd /etc/haproxy
```

**Download [haproxy.cfg file](#)**

**backup haproxy file**  
cp haproxy.cfg haproxy.cfg.bk

Replace the haproxy.cfg file with the one in Openshift4x-Installation folder  
vim haproxy.cfg

Add the following configuration lines to /etc/haproxy/haproxy.cfg:  
*(replace the IP addresses!)*

```
maxconn          3000

frontend openshift-api-server
    bind *:6443
    use_backend openshift-api-server
    mode tcp
    option tcplog

backend openshift-api-server
    balance source
    mode tcp
    server bootstrap-0 172.30.50.70:6443 check
    server control-plane-0 172.30.50.71:6443 check
    server control-plane-1 172.30.50.72:6443 check
    server control-plane-2 172.30.50.73:6443 check

frontend machine-config-server
    bind *:22623
    default_backend machine-config-server
    mode tcp
    option tcplog

backend machine-config-server
    balance source
    mode tcp
    server bootstrap-0 172.30.50.70:22623 check
    server control-plane-0 172.30.50.71:22623 check
    server control-plane-1 172.30.50.72:22623 check
    server control-plane-2 172.30.50.73:22623 check

frontend ingress-http
    bind *:80
    default_backend ingress-http
    mode tcp
    option tcplog
```

```
backend ingress-https
    balance source
    mode tcp
    server compute-0 172.30.50.74:443 check
    server compute-1 172.30.50.75:443 check
    server compute-2 172.30.50.76:443 check
    server compute-3 172.30.50.77:443 check
    server compute-4 172.30.50.78:443 check
```

Reboot server

## Log Back into Bastion node

Create core user

```
useradd core  
passwd core  
usermod -aG wheel core
```

Switch to core user

```
su - core
```

### Downloading required files

```
mkdir install_ocp4  
cd install_ocp4  
Install terraform version 0.11.14  
wget https://releases.hashicorp.com/terraform/0.11.14/terraform\_0.11.14\_linux\_amd64.zip  
unzip terraform_0.11.14_linux_amd64.zip  
sudo cp terraform /usr/local/bin  
terraform -v  
Will return the version of terraform  
Terraform v0.11.14
```

copy terraform installer form the downloaded repo

```
sudo cp /root/Openshift4x-Installation/vsphere -R /home/core/install_ocp4
```

Download ocp installer and client

```
wget https://mirror.openshift.com/pub/openshift-v4/clients/ocp/4.3.1/openshift-install-linux-4.3.1.tar.gz
```

```
wget https://mirror.openshift.com/pub/openshift-v4/clients/ocp/4.3.1/openshift-client-linux-4.3.1.tar.gz
```

```
tar -xvf openshift-install-linux-4.3.1.tar.gz  
tar -xvf openshift-client-linux-4.3.1.tar.gz
```

Download OVA image and upload to vSphere and convert OVA image to template

<https://mirror.openshift.com/pub/openshift-v4/dependencies/rhcos/>

Select the version you want to install

Since our CoreOS based VMs will be configured automatically, we have to provide a ssh public key to be able to log on to the nodes via ssh as user core  
ssh-keygen -t rsa -b 4096 -N "

```
start ssh-agent  
eval "$(ssh-agent -s)"
```

add your ssh private key to the ssh-agent  
ssh-add ~/.ssh/id\_rsa

For installing Red Hat OpenShift Container Platform 4x, you'll need a pull secret:

Link: <https://cloud.redhat.com/openshift/install/vsphere/user-provisioned>

copy your pull secret and save it in a text file

```
copy install-config.yaml from Openshift4x-Installation folder  
cp install-config.yaml install-config.yaml.bk  
vim install-config.yaml
```

**enter your vcenter details, pull secret and ssh-key.**

```
apiVersion: v1  
baseDomain: clientcenter.local  
compute:  
- hyperthreading: Enabled  
  name: worker  
  replicas: 5  
controlPlane:  
  hyperthreading: Enabled  
  name: master  
  replicas: 3  
metadata:  
  name: envisage  
networking:  
  clusterNetworks:  
  - cidr: 10.254.0.0/16  
    hostPrefix: 24  
  networkType: OpenShiftSDN  
  serviceNetwork:  
  - 192.168.0.0/16  
platform:  
vSphere:  
  vcenter: labvcsa.clientcentre.local  
  username: administrator@vsphere.local  
  password:   
  datacenter: Datacenter01  
  defaultDatastore: DsVol01  
pullSecret: '{"auths": {"cloud.openshift.com": {"auth": "b3B1bnNoaWZ0LXJ1bGVhc2UtZGV2K2FkYW1NEFHMTUyME9RTVBXQjVEU0kzV0c3RjJMSktJVzNXNTdTVTYxQTI=", "email": "adam.muganwa@ibm.com"}, "quWZiY2txaH1kc2lo0kMzRFcEyR1YxUTJYRUI2WFE3U1IyNk4yQzRTNEFHMTUyME9RTVBXQjVEU0kzV0c3RjJMSktJVzI3NDk2Nzd8dWhjLTFUcG5VbGVrZTdyRHZJYn1mYkNrUUhZRFNpaDpleUpoYkdjaU9pS1NvElV4TWLK0S5leUp6ZFd d2d4YmRHaG5YNkF0c2lMYnAtNjFBOWtjNDRHNURRVWhETWxxajBybWVvS1l3LWZLWkcxMlYwX2dzS1U1Vjh1MHB1Rg0NGFkS3FVbmZXWUx0TFRoTrNoLWVwcTB3aUVu0VdCUTVTZTJ5eHdHU3R3bE9i cGJsd2pxT215c21WT1psNTNael9a1lPVS12NUhXcWZpVHowUXVBWgtZNkRBR2NSZzUwZjI1cm1fSkFsb3h3d1BTcXBGb194RXhGcXQ10DBPRThZNDRjYBPd0poZ2d2NFJ6bTRfcGpaNXhXMzV1dXZES3BPCUV2dnp1d0JtMVNOukdZRmJkQ2Uz0WZvW1FtaXZIZ1FNQ3V1VU1NkNORkxucmFaMTI5eXQzSjdbdjBaWmFkdVRhRTY4dC1Rb3lmWU42X2lacj12UGRBaXJ4VVE4cm5TXzU4emxES3hTU==", "email": "adam.muganwa@ibm.com"}, "registry.redhat.io": {"auth": "NTI3NDk2Nzd8dWhjLTFUcG5QmhPRFUwWWpGak9EZzFaVEEyT1Rka05tWm10MkUxT0NKOS5UowlKQUFPNnFxN1FEMWdCd2d4YmRHaG5YNkF0c2lMYRYMG13Q1JIRnU0WGNVT2FielRvUHRubzBLR2QwVzMxaWtyZ1pnVlBjS1B3czJGa3BBcHg0NGFkS3FVbmZXWUx0TFRVVdONVVpWjUxdkhxaWtlbWJNbnc2aElhS1l0Wnp5NGx2M05maFJKVWtYN2xIblMxXz1ma1lPVS12NUhXcWZpVHowL1meFZKSXRnS3LERzY4b2V2ekpzSE4yQ1g5a0tRNu5ZMFb0WVrNxplbjZqZ0tRUUR3ZFBPd0poZ2d2NFJ6bTRfcGpcXZDZkc0UkZjve81WTNPuTlfZUZ0bkVmaUJSWUtQbGhZd1paOe02a0kwZGlwUVvzOG51NkNORkxucmFaMTI5eXQzSJpRz15NHdNV2NHdE1aekV0cTZCR3VQMTVzNUZHM1d1TVpQbmphZVh5bFdERUJWOC1kRQ==", "email": "adam.mug  
sshKey: ssh-rsa AAAAB3NzaC1yc2EAAAQABAAQACQCiLEjFQ8gsYzYGA1A/6hH6zAe5P81fKrZPJ12LTo8Wg@@@  
-- INSERT --
```

```
create ignition configs  
.openshift-install create ignition-configs
```

```
[core@Bastionhost install_ocp4]$ ./openshift-install create ignition-configs  
INFO Consuming "Install Config" from target directory
```

**ignition files are valid for 24 hours - so if your installation takes longer than 24 hours due to issues, you have to generate new ignition files.**

## Copy ignition files to your HTTP server

Copy the generated bootstrap.ign file to your HTTP server and ensure, that the file can be downloaded with http:

```
scp bootstrap.ign root@172.30.50.110:/var/www/html/
```

## Login to Load Balancer

```
chown apache.apache /var/www/html/bootstrap.ign
```

```
chmod 666 /var/www/html/bootstrap.ign
```

```
systemctl restart httpd
```

```
systemctl restart haproxy
```

### On Bastion

Check , if download would succeed from your http server.

```
curl -I http://172.30.50.110:8080/bootstrap.ign
[core@Bastionhost install_ocp4]$ curl -I http://172.30.50.110:8080/bootstrap.ign
HTTP/1.1 200 OK
Date: Thu, 30 Apr 2020 13:58:20 GMT
Server: Apache/2.4.37 (centos)
Last-Modified: Thu, 30 Apr 2020 13:38:13 GMT
ETag: "474f8-5a4822d27c36e"
Accept-Ranges: bytes
Content-Length: 292088
Content-Type: application/vnd.coreos.ignition+json
```

edit the terraform installer

```
cd /vsphere
```

```
vim terraform.tfvars
```

In sections control\_plane\_ignition / END\_OF\_MASTER\_IGNITION and compute\_ignition / END\_OF\_WORKER\_IGNITION,  
insert / copy&paste the contents of the ignitions files (master and worker) we generated before:  
END\_OF\_MASTER\_IGNITION -> master.ign  
END\_OF\_WORKER\_IGNITION -> worker.ign

```

cluster_id = "envisage"
cluster_domain = "envisage.clientcenter.local"
base_domain = "clientcenter.local"
vsphere_server = "labvcsa.clientcentre.local"
vsphere_user = "administrator@vsphere.local"
vsphere_password = ""
vsphere_cluster = "Cluster01"
vsphere_datacenter = "Datacenter01"
vsphere_datastore = "DsVol01"
vm_template = "rhcos-4.2.0"
machine_cidr = "172.30.50.0/24"

vm_network = "VM Network"

control_plane_count = 3
compute_count = 5

bootstrap_ignition_url = "http://172.30.50.110:8080/bootstrap.ign"

control_plane_ignition = <<END_OF_MASTER_IGNITION
{"ignition": {"config": {"append": [{"source": "https://api-int.envisage.clientcenter.local:22623/config/master", "verification": {}}], "security": .....}}
END_OF_MASTER_IGNITION

compute_ignition = <<END_OF_WORKER_IGNITION
{"ignition": {"config": {"append": [{"source": "https://api-int.envisage.clientcenter.local:22623/config/worker", "verification": {}}], "security": {"tls": .....}}}
END_OF_WORKER_IGNITION

# put your static IPs in here
bootstrap_ip = "172.30.50.70"
control_plane_ips = ["172.30.50.71", "172.30.50.72", "172.30.50.73"]
compute_ips = ["172.30.50.74", "172.30.50.75", "172.30.50.76", "172.30.50.77", "172.30.50.78"]

~
~

```

**Modify the file main.tf and delete or comment out the module "dns" part, since we do not want to use Route53/AWS as DNS provider:**

```

//module "dns" {
//  source = "./route53"
//
//  base_domain      = "${var.base_domain}"
//  cluster_domain   = "${var.cluster_domain}"
//  bootstrap_count  = "${var.bootstrap_complete ? 0 : 1}"
//  bootstrap_ips    = ["${module.bootstrap.ip_addresses}"]
//  control_plane_count = "${var.control_plane_count}"
//  control_plane_ips = ["${module.control_plane.ip_addresses}"]
//  compute_count    = "${var.compute_count}"
//  compute_ips      = ["${module.compute.ip_addresses}"]
//}

```

101.1

Bot

**configure DNS and network gateway, if needed**

vim vsphere/machine/ignition.tf

```
content {
  content = <<EOF
TYPE=Ethernet
BOOTPROTO=none
NAME=ens192
DEVICE=ens192
ONBOOT=yes
IPADDR=${local.ip_addresses[count.index]}
PREFIX=${local.mask}
GATEWAY=${local.gw}
DOMAIN=${var.cluster_domain}
DNS1=172.30.12.50
DNS2=9.9.9.9
EOF
}
}
```

### Set disk size

```
vim vsphere/machine/main.tf
```

```
network_interface {
  network_id = "${data.vsphere_network.network.id}"
}

disk {
  label        = "disk0"
  size         = 150
  thin_provisioned = "${data.vsphere_virtual_machine.template.disks.0.thin_provisioned"
}
```

### In the vSphere installer directory

```
terraform init
```

```
Terraform has been successfully initialized!
```

```
You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.
```

```
If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.
```

```
[core@Bastionhost vsphere]$ 
```

```
terraform plan
```

```
+ module.folder.vsphere_folder.folder
  id: <computed>
  datacenter_id: "datacenter-2"
  path: "envisage"
  type: "vm"

+ module.resource_pool.vsphere_resource_pool.resource_pool
  id: <computed>
  cpu_expandable: "true"
  cpu_limit: "-1"
  cpu_reservation: "0"
  cpu_share_level: "normal"
  cpu_shares: <computed>
  memory_expandable: "true"
  memory_limit: "-1"
  memory_reservation: "0"
  memory_share_level: "normal"
  memory_shares: <computed>
  name: "envisage"
  parent_resource_pool_id: "resgroup-8"
```

Plan: 11 to add, 0 to change, 0 to destroy.

---

Note: You didn't specify an "-out" parameter to save this plan, so Terraform can't guarantee that exactly these actions will be performed if "terraform apply" is subsequently run.

terraform apply -auto-approve

You should now see VMs in your vCenter



```

module.compute.vsphere_virtual_machine.vm[1]: Creation complete after 19s (ID: 4222e771-4
module.bootstrap.vsphere_virtual_machine.vm: Creation complete after 19s (ID: 42227317-c1
module.compute.vsphere_virtual_machine.vm[4]: Creation complete after 19s (ID: 42222298-b
module.compute.vsphere_virtual_machine.vm.3: Still creating... (20s elapsed)
module.control_plane.vsphere_virtual_machine.vm.0: Still creating... (20s elapsed)
module.control_plane.vsphere_virtual_machine.vm.2: Still creating... (20s elapsed)
module.compute.vsphere_virtual_machine.vm.0: Still creating... (20s elapsed)
module.compute.vsphere_virtual_machine.vm[2]: Creation complete after 20s (ID: 4222bad9-5
module.compute.vsphere_virtual_machine.vm[0]: Creation complete after 20s (ID: 4222a1f8-8
module.control_plane.vsphere_virtual_machine.vm[0]: Creation complete after 21s (ID: 4222
module.compute.vsphere_virtual_machine.vm[3]: Creation complete after 22s (ID: 42224f56-7
module.control_plane.vsphere_virtual_machine.vm[2]: Creation complete after 23s (ID: 4222

Apply complete! Resources: 11 added, 0 changed, 0 destroyed.
[core@bastion-prod vsphere]$ 

```

cd install\_ocp4 folder

./openshift-install --dir=. wait-for bootstrap-complete --log-level debug

```

[core@bastion-prod ocpinstall]$ openshift-install --dir=. wait-for bootstrap-complete --log-
DEBUG OpenShift Installer v4.3.5
DEBUG Built from commit 82f9a63c06956b3700a69475fbd14521e139aa1e
INFO Waiting up to 30m0s for the Kubernetes API at https://api.envi
DEBUG Still waiting for the Kubernetes API: the server could not find the requested resource
DEBUG Still waiting for the Kubernetes API: Get https://api.envi
DEBUG Still waiting for the Kubernetes API: Get https://api.envi
INFO API v1.16.2 up
INFO Waiting up to 30m0s for bootstrapping to complete...

```

If you want to see more details about installers progress, you can start a tail  
**tail -f openshift\_install.log**

cd install\_ocp4/vsphere

terraform apply -auto-approve -var 'bootstrap\_complete=true'  
cd install\_ocp4

```

vsphere_resource_pool.resource_pool: Refreshing state... (ID: resgroup-787)
vsphere_virtual_machine.vm: Refreshing state... (ID: 42296b34-fd02-9cbd-9ebf-c81c7dec5fbe)
vsphere_virtual_machine.vm[1]: Refreshing state... (ID: 422903b3-2b94-7bb2-3c3a-ff804b055a6b)
vsphere_virtual_machine.vm[0]: Refreshing state... (ID: 422910fb-e10d-abb0-93f3-2eca247c0a8d)
vsphere_virtual_machine.vm[2]: Refreshing state... (ID: 4229c635-4a40-f4e2-cb27-a759682db6d1)
vsphere_virtual_machine.vm[0]: Refreshing state... (ID: 422919d6-cb13-bd82-361f-ec2c7bd1c642)
vsphere_virtual_machine.vm[1]: Refreshing state... (ID: 42294e9d-84d5-56e8-1b73-d80bf9c5d6f8)
vsphere_virtual_machine.vm[2]: Refreshing state... (ID: 42290a2d-6c83-398f-ce53-732a12f2057e)
vsphere_virtual_machine.vm[3]: Refreshing state... (ID: 42293da0-4ff9-1b02-92d3-2c500e53f788)
vsphere_virtual_machine.vm[4]: Refreshing state... (ID: 422946d0-227b-8858-4c9a-71dfc568f4af)
module.bootstrap.vsphere_virtual_machine.vm: Destroying... (ID: 42296b34-fd02-9cbd-9ebf-c81c7de
module.bootstrap.vsphere_virtual_machine.vm: Destruction complete after 10s

Apply complete! Resources: 0 added, 0 changed, 1 destroyed.
[core@Bastionhost vsphere]$ 

```

./openshift-install --dir=. wait-for install-complete

```
[core@bastion-prod ocpinstall]$ ./openshift-install --dir=
INFO Waiting up to 30m0s for the cluster at https://api.en
INFO Waiting up to 10m0s for the openshift-console route t
INFO Install complete!
INFO To access the cluster as the system:admin user when u
INFO Access the OpenShift web-console here: https://consol
INFO Login to the console with user: kubeadmin, password:
[core@bastion-prod ocpinstall]$
```